

# Lecture 3: SVM dual, kernels and regression

---

C19 Machine Learning

Hilary 2015

A. Zisserman

---

- Primal and dual forms
- Linear separability revisited
- Feature maps
- Kernels for SVMs
- Regression
  - Ridge regression
  - Basis functions

# SVM – review

---

- We have seen that for an SVM learning a linear classifier

$$f(x) = \mathbf{w}^\top \mathbf{x} + b$$

is formulated as solving an optimization problem over  $\mathbf{w}$  :

$$\min_{\mathbf{w} \in \mathbb{R}^d} \|\mathbf{w}\|^2 + C \sum_i^N \max(0, 1 - y_i f(\mathbf{x}_i))$$

- This quadratic optimization problem is known as the **primal** problem.

- Instead, the SVM can be formulated to learn a linear classifier

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i (\mathbf{x}_i^\top \mathbf{x}) + b$$

by solving an optimization problem over  $\alpha_i$ .

- This is known as the **dual** problem, and we will look at the advantages of this formulation.

# Sketch derivation of dual form

---

The [Representer Theorem](#) states that the solution  $\mathbf{w}$  can always be written as a linear combination of the training data:

$$\mathbf{w} = \sum_{j=1}^N \alpha_j y_j \mathbf{x}_j$$

Proof: [see example sheet](#) .

Now, substitute for  $\mathbf{w}$  in  $f(x) = \mathbf{w}^\top \mathbf{x} + b$

$$f(x) = \left( \sum_{j=1}^N \alpha_j y_j \mathbf{x}_j \right)^\top \mathbf{x} + b = \sum_{j=1}^N \alpha_j y_j (\mathbf{x}_j^\top \mathbf{x}) + b$$

and for  $\mathbf{w}$  in the cost function  $\min_{\mathbf{w}} \|\mathbf{w}\|^2$  subject to  $y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \forall i$

$$\|\mathbf{w}\|^2 = \left\{ \sum_j \alpha_j y_j \mathbf{x}_j \right\}^\top \left\{ \sum_k \alpha_k y_k \mathbf{x}_k \right\} = \sum_{jk} \alpha_j \alpha_k y_j y_k (\mathbf{x}_j^\top \mathbf{x}_k)$$

Hence, an equivalent optimization problem is over  $\alpha_j$

$$\min_{\alpha_j} \sum_{jk} \alpha_j \alpha_k y_j y_k (\mathbf{x}_j^\top \mathbf{x}_k) \quad \text{subject to} \quad y_i \left( \sum_{j=1}^N \alpha_j y_j (\mathbf{x}_j^\top \mathbf{x}_i) + b \right) \geq 1, \forall i$$

and a few more steps are required to complete the derivation.

# Primal and dual formulations

---

$N$  is number of training points, and  $d$  is dimension of feature vector  $\mathbf{x}$ .

Primal problem: for  $\mathbf{w} \in \mathbb{R}^d$

$$\min_{\mathbf{w} \in \mathbb{R}^d} \|\mathbf{w}\|^2 + C \sum_i^N \max(0, 1 - y_i f(\mathbf{x}_i))$$

Dual problem: for  $\alpha \in \mathbb{R}^N$  (stated without proof):

$$\max_{\alpha_i \geq 0} \sum_i \alpha_i - \frac{1}{2} \sum_{j,k} \alpha_j \alpha_k y_j y_k (\mathbf{x}_j^\top \mathbf{x}_k) \text{ subject to } 0 \leq \alpha_i \leq C \text{ for } \forall i, \text{ and } \sum_i \alpha_i y_i = 0$$

- Need to learn  $d$  parameters for primal, and  $N$  for dual
- If  $N \ll d$  then more efficient to solve for  $\alpha$  than  $\mathbf{w}$
- Dual form only involves  $(\mathbf{x}_j^\top \mathbf{x}_k)$ . We will return to why this is an advantage when we look at kernels.

# Primal and dual formulations

---

Primal version of classifier:


$$f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$$

Dual version of classifier:


$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i (\mathbf{x}_i^\top \mathbf{x}) + b$$

At first sight the dual form appears to have the disadvantage of a K-NN classifier – it requires the training data points  $\mathbf{x}_i$ . However, many of the  $\alpha_i$ 's are zero. The ones that are non-zero define the support vectors  $\mathbf{x}_i$ .

# Support Vector Machine



$C = 10$  soft margin




Comment Window

SVM (L1) by Sequential Minimal Optimizer  
Kernel: linear (-), C: 10.0000  
Kernel evaluations: 2645  
Number of Support Vectors: 4  
Margin: 0.2265  
Training error: 3.70%

# Handling data that is not linearly separable

---




- introduce slack variables

$$\min_{\mathbf{w} \in \mathbb{R}^d, \xi_i \in \mathbb{R}^+} \|\mathbf{w}\|^2 + C \sum_i^N \xi_i$$

subject to

$$y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i \text{ for } i = 1 \dots N$$




- linear classifier not appropriate

??



# Solution 1: use polar coordinates

---




- Data **is** linearly separable in polar coordinates
- Acts non-linearly in original space

$$\Phi : \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} r \\ \theta \end{pmatrix} \quad \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

## Solution 2: map data to higher dimension

---


$$\Phi : \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{pmatrix} \quad \mathbb{R}^2 \rightarrow \mathbb{R}^3$$



- Data **is** linearly separable in 3D
- This means that the problem can still be solved by a linear classifier

# SVM classifiers in a transformed feature space

---



Learn classifier linear in  $\mathbf{w}$  for  $\mathbb{R}^D$ :

$$f(\mathbf{x}) = \mathbf{w}^\top \Phi(\mathbf{x}) + b$$

$\Phi(\mathbf{x})$  is a **feature map**

# Primal Classifier in transformed feature space

---

Classifier, with  $\mathbf{w} \in \mathbb{R}^D$ :

$$f(\mathbf{x}) = \mathbf{w}^\top \Phi(\mathbf{x}) + b$$

Learning, for  $\mathbf{w} \in \mathbb{R}^D$

$$\min_{\mathbf{w} \in \mathbb{R}^D} \|\mathbf{w}\|^2 + C \sum_i^N \max(0, 1 - y_i f(\mathbf{x}_i))$$

- Simply map  $\mathbf{x}$  to  $\Phi(\mathbf{x})$  where data is separable
- Solve for  $\mathbf{w}$  in high dimensional space  $\mathbb{R}^D$
- If  $D \gg d$  then there are many more parameters to learn for  $\mathbf{w}$ . Can this be avoided?

# Dual Classifier in transformed feature space

---

Classifier:

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i \mathbf{x}_i^\top \mathbf{x} + b$$
$$\rightarrow f(\mathbf{x}) = \sum_i^N \alpha_i y_i \Phi(\mathbf{x}_i)^\top \Phi(\mathbf{x}) + b$$

Learning:

$$\max_{\alpha_i \geq 0} \sum_i \alpha_i - \frac{1}{2} \sum_{jk} \alpha_j \alpha_k y_j y_k \mathbf{x}_j^\top \mathbf{x}_k$$
$$\rightarrow \max_{\alpha_i \geq 0} \sum_i \alpha_i - \frac{1}{2} \sum_{jk} \alpha_j \alpha_k y_j y_k \Phi(\mathbf{x}_j)^\top \Phi(\mathbf{x}_k)$$

subject to

$$0 \leq \alpha_i \leq C \text{ for } \forall i, \text{ and } \sum_i \alpha_i y_i = 0$$

# Dual Classifier in transformed feature space

---

- Note, that  $\Phi(\mathbf{x})$  only occurs in pairs  $\Phi(\mathbf{x}_j)^\top \Phi(\mathbf{x}_i)$
- Once the scalar products are computed, only the  $N$  dimensional vector  $\alpha$  needs to be learnt; it is not necessary to learn in the  $D$  dimensional space, as it is for the primal
- Write  $k(\mathbf{x}_j, \mathbf{x}_i) = \Phi(\mathbf{x}_j)^\top \Phi(\mathbf{x}_i)$ . This is known as a **Kernel**

Classifier:

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) + b$$

Learning:

$$\max_{\alpha_i \geq 0} \sum_i \alpha_i - \frac{1}{2} \sum_{jk} \alpha_j \alpha_k y_j y_k k(\mathbf{x}_j, \mathbf{x}_k)$$

subject to

$$0 \leq \alpha_i \leq C \text{ for } \forall i, \text{ and } \sum_i \alpha_i y_i = 0$$

# Special transformations

---

$$\Phi : \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} x_1^2 \\ x_2^2 \\ \sqrt{2}x_1x_2 \end{pmatrix} \quad \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$\begin{aligned} \Phi(\mathbf{x})^\top \Phi(\mathbf{z}) &= (x_1^2, x_2^2, \sqrt{2}x_1x_2) \begin{pmatrix} z_1^2 \\ z_2^2 \\ \sqrt{2}z_1z_2 \end{pmatrix} \\ &= x_1^2z_1^2 + x_2^2z_2^2 + 2x_1x_2z_1z_2 \\ &= (x_1z_1 + x_2z_2)^2 \\ &= (\mathbf{x}^\top \mathbf{z})^2 \end{aligned}$$

## Kernel Trick

- Classifier can be **learnt** and **applied** without explicitly computing  $\Phi(\mathbf{x})$
- All that is required is the kernel  $k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^\top \mathbf{z})^2$
- Complexity of learning depends on  $N$  (typically it is  $O(N^3)$ ) not on  $D$

# Example kernels

---

- **Linear** kernels  $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{x}'$
- **Polynomial** kernels  $k(\mathbf{x}, \mathbf{x}') = (1 + \mathbf{x}^\top \mathbf{x}')^d$  for any  $d > 0$ 
  - Contains all polynomials terms up to degree  $d$
- **Gaussian** kernels  $k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2 / 2\sigma^2)$  for  $\sigma > 0$ 
  - Infinite dimensional feature space



## SVM classifier with Gaussian kernel

---

$N$  = size of training data

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}) + b$$

weight (may be zero)

support vector


Gaussian kernel  $k(\mathbf{x}, \mathbf{x}') = \exp(-||\mathbf{x} - \mathbf{x}'||^2 / 2\sigma^2)$

### Radial Basis Function (RBF) SVM

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i \exp(-||\mathbf{x} - \mathbf{x}_i||^2 / 2\sigma^2) + b$$

# RBF Kernel SVM Example


---



- data is not linearly separable in original feature space

$$\sigma = 1.0 \quad C = \infty$$

$$f(\mathbf{x}) = 1$$



SMO (L1)

Kernel

RBF

Kernel argument

1

C-constant

Inf

epsilon,tolerance

1e-3,1e-3

☒ Background

Comment Window

SVM (L1) by Sequential Minimal Optimizer

Kernel: rbf (1), C: Inf

Kernel evaluations: 321750


Number of Support Vectors: 5

Margin: 0.0440

Training error: 0.00%


$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i \exp \left( -\|\mathbf{x} - \mathbf{x}_i\|^2 / 2\sigma^2 \right) + b$$

$$\sigma = 1.0 \quad C = 100$$



Decrease C, gives wider (soft) margin

$$\sigma = 1.0 \quad C = 10$$



SMO (L1)

Kernel

RBF

Kernel argument

1

C-constant

10

epsilon,tolerance

1e-3,1e-3

☒ Background

Comment Window

SVM (L1) by Sequential Minimal Optimizer

Kernel: rbf (1), C: 10.0000

Kernel evaluations: 46158


Number of Support Vectors: 24

Margin: 0.0755

Training error: 0.00%

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i \exp \left( -\|\mathbf{x} - \mathbf{x}_i\|^2 / 2\sigma^2 \right) + b$$

$$\sigma = 1.0 \quad C = \infty$$



SMO (L1) ▼

Kernel

RBF ▼

Kernel argument

1

C-constant

Inf

epsilon,tolerance

1e-3,1e-3

☒ Background

Load data

Create data

Reset

Train SVM

Info

Close

Comment Window

SVM (L1) by Sequential Minimal Optimizer

Kernel: rbf (1), C: Inf

Kernel evaluations: 62739


Number of Support Vectors: 5

Margin: 0.0445

Training error: 0.00%

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i \exp \left( -\|\mathbf{x} - \mathbf{x}_i\|^2 / 2\sigma^2 \right) + b$$

$$\sigma = 0.25 \quad C = \infty$$



SMO (L1)

Kernel

RBF

Kernel argument

0.25

C-constant

Inf

epsilon,tolerance

1e-3,1e-3

☒ Background

Comment Window

SVM (L1) by Sequential Minimal Optimizer

Kernel: rbf (0.25), C: Inf

Kernel evaluations: 42795


Number of Support Vectors: 18

Margin: 0.2358

Training error: 0.00%

Decrease sigma, moves towards nearest neighbour classifier

$$\sigma = 0.1 \quad C = \infty$$



SMO (L1)

Kernel

RBF

Kernel argument

0.1

C-constant

Inf

epsilon,tolerance

1e-3,1e-3

☒ Background

Comment Window

SVM (L1) by Sequential Minimal Optimizer

Kernel: rbf (0.1), C: Inf

Kernel evaluations: 173935

Number of Support Vectors: 62

Margin: 0.2196

Training error: 0.00%

$$f(\mathbf{x}) = \sum_i^N \alpha_i y_i \exp \left( -\|\mathbf{x} - \mathbf{x}_i\|^2 / 2\sigma^2 \right) + b$$



# Kernel Trick - Summary


---

- Classifiers can be learnt for high dimensional features spaces, without actually having to map the points into the high dimensional space
- Data may be linearly separable in the high dimensional space, but not linearly separable in the original feature space
- Kernels can be used for an SVM because of the scalar product in the dual form, but can also be used elsewhere – they are not tied to the SVM formalism
- Kernels apply also to objects that are not vectors, e.g.

$$k(h, h') = \sum_k \min(h_k, h'_k) \text{ for histograms with bins } h_k, h'_k$$

# Regression

---



- Suppose we are given a training set of  $N$  observations

$$((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)) \text{ with } \mathbf{x}_i \in \mathbb{R}^d, y_i \in \mathbb{R}$$

- The **regression** problem is to estimate  $f(\mathbf{x})$  from this data such that

$$y_i = f(\mathbf{x}_i)$$

# Learning by optimization

---

- As in the case of classification, learning a regressor can be formulated as an optimization:

Minimize with respect to  $f \in \mathcal{F}$

$$\sum_{i=1}^N \underbrace{l(f(\mathbf{x}_i), y_i)}_{\text{loss function}} + \underbrace{\lambda R(f)}_{\text{regularization}}$$

- There is a choice of both loss functions and regularization
  - e.g. squared loss, SVM “hinge-like” loss
  - squared regularizer, lasso regularizer

## Choice of regression function – non-linear basis functions

---

- Function for regression  $y(\mathbf{x}, \mathbf{w})$  is a **non-linear function** of  $\mathbf{x}$ , but **linear** in  $\mathbf{w}$ :

$$f(\mathbf{x}, \mathbf{w}) = w_0 + w_1\phi_1(\mathbf{x}) + w_2\phi_2(\mathbf{x}) + \dots + w_M\phi_M(\mathbf{x}) = \mathbf{w}^\top \Phi(\mathbf{x})$$

- For example, for  $x \in \mathbb{R}$ , polynomial regression with  $\phi_j(x) = x^j$ :

$$f(x, \mathbf{w}) = w_0 + w_1\phi_1(\mathbf{x}) + w_2\phi_2(\mathbf{x}) + \dots + w_M\phi_M(\mathbf{x}) = \sum_{j=0}^M w_j x^j$$

e.g. for  $M = 3$ ,

$$f(x, \mathbf{w}) = (w_0, w_1, w_2, w_3) \begin{pmatrix} 1 \\ x \\ x^2 \\ x^3 \end{pmatrix} = \mathbf{w}^\top \Phi(x)$$


$$\Phi : x \rightarrow \Phi(x) \quad \mathbb{R}^1 \rightarrow \mathbb{R}^4$$

# Least squares “ridge regression”

- Cost function – squared loss:

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \underbrace{\{f(x_i, \mathbf{w}) - y_i\}^2}_{\text{loss function}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{regularization}}$$

target value



- Regression function for x (1D):

$$f(\mathbf{x}, \mathbf{w}) = w_0 + w_1 \phi_1(\mathbf{x}) + w_2 \phi_2(\mathbf{x}) + \dots + w_M \phi_M(\mathbf{x}) = \mathbf{w}^\top \Phi(\mathbf{x})$$

- NB squared loss arises in Maximum Likelihood estimation for an error model

$$y_i = \tilde{y}_i + n_i \quad n_i \sim \mathcal{N}(0, \sigma^2)$$

measured value      true value

# Solving for the weights $\mathbf{w}$

---

Notation: write the target and regressed values as  $N$ -vectors

$$\mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix} \quad \mathbf{f} = \begin{pmatrix} \Phi(x_1)^\top \mathbf{w} \\ \Phi(x_2)^\top \mathbf{w} \\ \vdots \\ \Phi(x_N)^\top \mathbf{w} \end{pmatrix} = \Phi \mathbf{w} = \begin{bmatrix} 1 & \phi_1(x_1) & \dots & \phi_M(x_1) \\ 1 & \phi_1(x_2) & \dots & \phi_M(x_2) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \phi_1(x_N) & \dots & \phi_M(x_N) \end{bmatrix} \begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_M \end{pmatrix}$$

$\Phi$  is an  $N \times M$  **design matrix**

e.g. for polynomial regression with basis functions up to  $x^2$

$$\Phi \mathbf{w} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 \end{bmatrix} \begin{pmatrix} w_0 \\ w_1 \\ w_2 \end{pmatrix}$$

$$\begin{aligned}
\tilde{E}(\mathbf{w}) &= \frac{1}{2} \sum_{i=1}^N \{f(x_i, \mathbf{w}) - y_i\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2 \\
&= \frac{1}{2} \sum_{i=1}^N \left(y_i - \mathbf{w}^\top \Phi(x_i)\right)^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2 \\
&= \frac{1}{2} (\mathbf{y} - \Phi \mathbf{w})^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2
\end{aligned}$$

Now, compute where derivative w.r.t.  $\mathbf{w}$  is zero for minimum

$$\frac{\tilde{E}(\mathbf{w})}{d\mathbf{w}} = -\Phi^\top (\mathbf{y} - \Phi \mathbf{w}) + \lambda \mathbf{w} = 0$$

Hence

$$(\Phi^\top \Phi + \lambda \mathbf{I}) \mathbf{w} = \Phi^\top \mathbf{y}$$

$$\mathbf{w} = (\Phi^\top \Phi + \lambda \mathbf{I})^{-1} \Phi^\top \mathbf{y}$$

M basis functions, N data points

$$\mathbf{w} = \left( \Phi^T \Phi + \lambda \mathbf{I} \right)^{-1} \Phi^T \mathbf{y}$$

$$\begin{matrix} \left( \begin{matrix} \end{matrix} \right) & = & \left( \begin{matrix} \end{matrix} \right) & \left( \begin{matrix} \end{matrix} \right) & \left( \begin{matrix} \end{matrix} \right) \end{matrix} \quad \text{assume } N > M$$

$M \times 1$ 
 $M \times M$ 
 $M \times N$ 
 $N \times 1$

- This shows that there is a unique solution.
- If  $\lambda = 0$  (no regularization), then

$$\mathbf{w} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y} = \Phi^+ \mathbf{y}$$

where  $\Phi^+$  is the [pseudo-inverse](#) of  $\Phi$  (`pinv` in Matlab)

- Adding the term  $\lambda \mathbf{I}$  improves the [conditioning](#) of the inverse, since if  $\Phi$  is not full rank, then  $(\Phi^T \Phi + \lambda \mathbf{I})$  will be (for sufficiently large  $\lambda$ )
- As  $\lambda \rightarrow \infty$ ,  $\mathbf{w} \rightarrow \frac{1}{\lambda} \Phi^T \mathbf{y} \rightarrow \mathbf{0}$
- Often the regularization is applied only to the inhomogeneous part of  $\mathbf{w}$ , i.e. to  $\tilde{\mathbf{w}}$ , where  $\mathbf{w} = (w_0, \tilde{\mathbf{w}})$



$$\mathbf{w} = \left( \Phi^\top \Phi + \lambda \mathbf{I} \right)^{-1} \Phi^\top \mathbf{y}$$

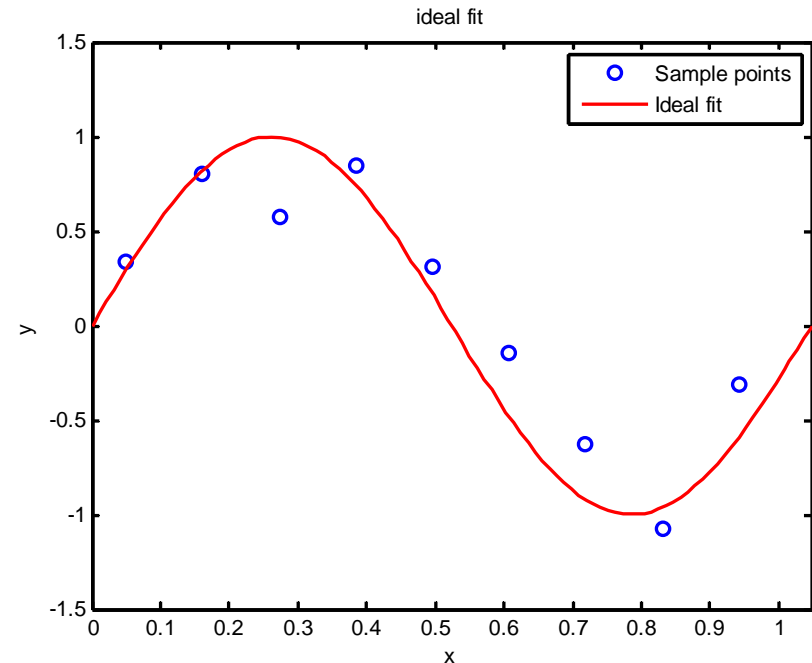
$$\begin{aligned} f(x, \mathbf{w}) &= \mathbf{w}^\top \Phi(x) = \Phi(x)^\top \mathbf{w} \\ &= \Phi(x)^\top \left( \Phi^\top \Phi + \lambda \mathbf{I} \right)^{-1} \Phi^\top \mathbf{y} \\ &= \mathbf{b}(x)^\top \mathbf{y} \end{aligned}$$

Output is a **linear blend**,  $\mathbf{b}(x)$ , of the training values  $\{y_i\}$

# Example 1: polynomial basis functions

---

- The red curve is the true function (which is not a polynomial)
- The data points are samples from the curve with added noise in y.
- There is a choice in both the degree,  $M$ , of the basis functions used, and in the strength of the regularization







$$f(x, \mathbf{w}) = \sum_{j=0}^M w_j x^j = \mathbf{w}^\top \Phi(x) \quad \Phi : x \rightarrow \Phi(x) \quad \mathbb{R} \rightarrow \mathbb{R}^{M+1}$$



$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \{f(x_i, \mathbf{w}) - y_i\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$\mathbf{w}$  is a  $M+1$  dimensional vector



$N = 9$  samples,  $M = 7$



$M = 3$




$M = 5$



## Example 2: Gaussian basis functions

---

- The red curve is the true function (which is not a polynomial)
- The data points are samples from the curve with added noise in  $y$ .
- Basis functions are centred on the training data ( $N$  points)
- There is a choice in both the scale, sigma, of the basis functions used, and in the strength of the regularization


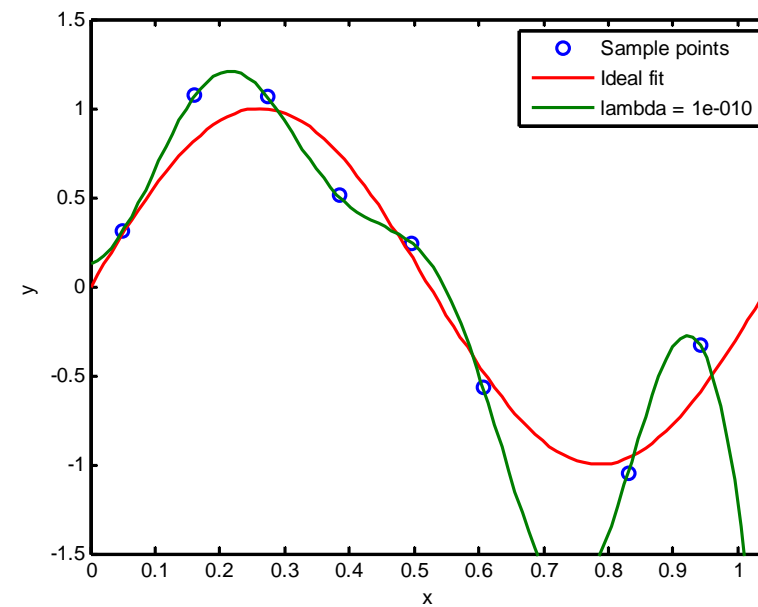




$$f(x, \mathbf{w}) = \sum_{i=1}^N w_i e^{-(x-x_i)^2/\sigma^2} = \mathbf{w}^\top \Phi(x) \quad \Phi : x \rightarrow \Phi(x) \quad \mathbb{R} \rightarrow \mathbb{R}^N$$

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \{f(x_i, \mathbf{w}) - y_i\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



$\mathbf{w}$  is a  $N$ -vector

$N = 9$  samples,  $\sigma = 0.334$




# Choosing lambda using a validation set


---




Sigma = 0.334




Sigma = 0.1



Gaussian basis functions



Gaussian basis functions






# Application: regressing face pose

---

- Estimate two face pose angles:
  - yaw (around the Y axis)
  - pitch (around the X axis)
- Compute a HOG feature vector for each face region
- Learn a regressor from the HOG vector to the two pose angles



# Summary and dual problem

---

So far we have considered the **primal** problem where

$$f(\mathbf{x}, \mathbf{w}) = \sum_{i=1}^M w_i \phi_i(\mathbf{x}) = \mathbf{w}^\top \Phi(\mathbf{x})$$

and we wanted a solution for  $\mathbf{w} \in \mathbb{R}^M$

As in the case of SVMs, we can also consider the **dual** problem where

$$\mathbf{w} = \sum_{i=1}^N \mathbf{a}_i \Phi(x_i) \quad \text{and} \quad f(\mathbf{x}, \mathbf{a}) = \sum_i a_i \Phi(x_i)^\top \Phi(x)$$

and obtain a solution for  $\mathbf{a} \in \mathbb{R}^N$ .

Again

- there is a closed form solution for  $\mathbf{a}$ ,
- the solution involves the  $N \times N$  Gram matrix  $k(x_i, x_j) = \Phi(x_i)^\top \Phi(x_j)$ ,
- so we can use the kernel trick again to replace scalar products

## Background reading and more

---

- Bishop, chapters 6 & 7 for kernels and SVMs
- Hastie et al, chapter 12
- Bishop, chapter 3 for regression
- More on web page:  
<http://www.robots.ox.ac.uk/~az/lectures/ml>